

Observability for On-Prem AI Gateways with Grafana

Ronald Rahaman

Research Scientist II

The Partnership for an Advanced Computing Environment

Georgia Institute of Technology Office of Information Technology

What are we doing?

We're deploying an on-prem AI Gateway to accompany [GT's AI MakerSpace \(AIMS\)](#)

An AI Gateway provides API access to LLMs for end-users' apps

- Gateway is middleware between client and inference engines
- Provides authentication, accounting, rate limiting, etc.

PACE's gateway (SPIN) is composed of FOSS with proven commercial usage

- vLLM: Inference engine
- LiteLLM: Rate limiting and accounting
- LiteMaaS: Authentication and authorization




What do we need?

- LiteLLM and LiteMaaS provide admin API endpoints for observability
 - Usage: Token consumption per user, team
 - Spend: User and Team budgets
 - Health checks: Availability of inference engines
- Goal: Hook admin API into Grafana dashboards for visual observability and automatic notifications
- Problem: Existing LiteLLM + Grafana integration doesn't provide the info and visuals we need



What will you do?



- Develop a Grafana Dashboard for on-prem LLMs that is:
 - Published on GitHub under an open source license
 - Listed in Grafana Labs' dashboards collection
 - Deployed at PACE
- Benefits for you!
 - Real-world experience in AIOps
 - Hone skills in API-first app design
 - A FOSS app used in IT operations for your resume